

```
In [1]: import torch
```

## 尽量不要自己给自己赋值

### 错误例子：自己给自己赋值

```
In [2]: etal = torch.nn.Parameter(torch.rand(10), requires_grad=True)
        etal

Out[2]: Parameter containing:
        tensor([0.0312, 0.6983, 0.4285, 0.4792, 0.2871, 0.7417, 0.1106, 0.9623, 0.9452,
        0.4848], requires_grad=True)
        这里本来想让 etal 乘以2

In [3]: etal = etal * 2

In [4]: etal

Out[4]: tensor([0.0625, 1.3966, 0.8570, 0.9584, 0.5741, 1.4835, 0.2212, 1.9247, 1.8905,
        0.9695], grad_fn=<MulBackward0>)

In [5]: etal.sum().backward()

结果可以看出没有梯度了

In [6]: etal.grad
```

```
/Users/haibinzhao/miniconda3/envs/ML/lib/python3.8/site-packages/torch/_tensor.py:1104: UserWarning: The .grad attribute of a Tensor that is not a leaf Tensor is being accessed. Its .grad attribute won't be populated during autograd.backward(). If you indeed want the .grad field to be populated for a non-leaf Tensor, use .retain_grad() on the non-leaf Tensor. If you access the non-leaf Tensor by mistake, make sure you access the leaf Tensor instead. See github.com/pytorch/pytorch/pull/30531 for more informations. (Triggered internally at /Users/distiller/project/conda/conda-bld/pytorch_1646755922314/work/build/aten/src/ATen/core/TensorBody.h:475.)
  return self._grad
```

### 正确例子：赋值给其他变量名

```
In [7]: eta2 = torch.nn.Parameter(torch.rand(10), requires_grad=True)
        eta2

Out[7]: Parameter containing:
        tensor([0.8627, 0.4348, 0.9095, 0.5124, 0.3881, 0.6686, 0.7270, 0.6214, 0.1460,
        0.6846], requires_grad=True)

In [8]: eta_temp = eta2 * 2

In [9]: eta_temp

Out[9]: tensor([1.7255, 0.8697, 1.8190, 1.0247, 0.7762, 1.3372, 1.4540, 1.2428, 0.2920,
        1.3692], grad_fn=<MulBackward0>)

In [10]: eta_temp.sum().backward()

In [11]: eta2.grad

Out[11]: tensor([2., 2., 2., 2., 2., 2., 2., 2., 2., 2.]
```

## 小心直接赋值其他变量

### 错误例子：直接赋值给别的变量后，内存里还是同一个变量

```
In [12]: v1 = torch.nn.Parameter(torch.rand(10), requires_grad=True)
        v1

Out[12]: Parameter containing:
        tensor([0.2928, 0.6194, 0.0654, 0.9038, 0.8469, 0.8558, 0.8001, 0.8354, 0.2006,
        0.9793], requires_grad=True)
        把 v1 赋值给 a

In [13]: a = v1[:4]

修改 a 的值

In [15]: a.data.copy_(torch.rand(4))
        a

Out[15]: tensor([0.1814, 0.2235, 0.1048, 0.1320], grad_fn=<SliceBackward0>)

发现 v1 的值也被改了

In [16]: v1

Out[16]: Parameter containing:
        tensor([0.1814, 0.2235, 0.1048, 0.1320, 0.8469, 0.8558, 0.8001, 0.8354, 0.2006,
        0.9795], requires_grad=True)
```

### 正确例子：用克隆的方法

```
In [17]: v2 = torch.nn.Parameter(torch.rand(10), requires_grad=True)
        v2

Out[17]: Parameter containing:
        tensor([0.6647, 0.6194, 0.3355, 0.2664, 0.4136, 0.1148, 0.5026, 0.8328, 0.2710,
        0.9257], requires_grad=True)

In [18]: b = v2[:4].clone()

In [19]: b.data.copy_(torch.rand(4))
        b

Out[19]: tensor([0.3222, 0.6089, 0.8604, 0.0538], grad_fn=<CloneBackward0>)

In [20]: v2

Out[20]: Parameter containing:
        tensor([0.6647, 0.6194, 0.3355, 0.2664, 0.4136, 0.1148, 0.5026, 0.8328, 0.2710,
        0.9257], requires_grad=True)
```

## 生成可学习参数时先生成，最后再包装

### 错误例子：先生成Parameter，再变形

```
In [21]: t = torch.nn.Parameter(torch.randn(12,1), requires_grad=True)
        t

Out[21]: Parameter containing:
        tensor([[-0.5475],
        [ 2.0614],
        [-0.0118],
        [ 1.6362],
        [ 0.0855],
        [ 0.2819],
        [-1.6987],
        [-1.1044],
        [ 1.1377],
        [ 0.4085],
        [ 1.4628],
        [-0.2824]], requires_grad=True)

In [23]: T = t.view(3,4)
        T

Out[23]: tensor([[[-0.5475,  2.0614, -0.0118,  1.6362],
        [ 0.0855,  0.2819, -1.6987, -1.1044],
        [ 1.1377,  0.4085,  1.4628, -0.2824]], grad_fn=<ViewBackward0>)

In [24]: T.sum().backward()

In [25]: T.grad

/Users/haibinzhao/miniconda3/envs/ML/lib/python3.8/site-packages/torch/_tensor.py:1104: UserWarning: The .grad attribute of a Tensor that is not a leaf Tensor is being accessed. Its .grad attribute won't be populated during autograd.backward(). If you indeed want the .grad field to be populated for a non-leaf Tensor, use .retain_grad() on the non-leaf Tensor. If you access the non-leaf Tensor by mistake, make sure you access the leaf Tensor instead. See github.com/pytorch/pytorch/pull/30531 for more informations. (Triggered internally at /Users/distiller/project/conda/conda-bld/pytorch_1646755922314/work/build/aten/src/ATen/core/TensorBody.h:475.)
  return self._grad
```

### 正确例子：先全部初始化完成，再包装成Parameter

```
In [24]: k = torch.randn(12,1)
        k

Out[24]: tensor([[-0.2693],
        [ 0.2612],
        [-0.8053],
        [ 0.5455],
        [ 1.5656],
        [-1.4120],
        [ 1.5922],
        [-1.7210],
        [ 0.8812],
        [ 0.9397],
        [-0.9860],
        [ 1.3318]])

In [25]: k = k.view(3,4)

In [26]: kp = torch.nn.Parameter(k, requires_grad=True)

In [27]: kp.sum().backward()

In [28]: kp.grad

Out[28]: tensor([[1., 1., 1., 1.],
        [1., 1., 1., 1.],
        [1., 1., 1., 1.]])
```

## 组装变量时的错误

### 错误例子：用tensor把变量重新拼起来

```
In [26]: a = torch.nn.Parameter(torch.rand(5), requires_grad=True)

In [27]: a0 = a[0] * 2
        a1 = torch.log(a[1]*1)
        a2 = a[2] ** 2
        a3 = torch.nn.functional.softplus(a[3] * 2)
        a4 = a[4]
        a0, a1, a2, a3, a4

Out[27]: (tensor(0.1184, grad_fn=<MulBackward0>),
        tensor(0.3061, grad_fn=<LogBackward0>),
        tensor(0.0003, grad_fn=<PowBackward0>),
        tensor(1.3874, grad_fn=<SoftplusBackward0>),
        tensor(0.5207, grad_fn=<SelectBackward0>))

组合成一个新的变量

In [28]: A1 = torch.tensor([a0, a1, a2, a3, a4])
        A1

Out[28]: tensor([1.1843e-01, 3.0607e-01, 2.5946e-04, 1.3874e+00, 5.2072e-01])

由于重新用 tensor 包装了变量，切断了反向传播

In [29]: A1.sum().backward()

RuntimeError                                Traceback (most recent call last)
Input In [29], in <cell line: 1>()
----> 1 A1.sum().backward()

File ~/miniconda3/envs/ML/lib/python3.8/site-packages/torch/_tensor.py:363, in Tensor.backward(self, gradient, retain_graph, create_graph, inputs)
    354 if has_torch_function_unary(self):
    355     return handle_torch_function(
    356         Tensor.backward,
    357         (self,),
    358         (...)
    361         create_graph=create_graph,
    362         inputs=inputs)
--> 363 torch.autograd.backward(self, gradient, retain_graph, create_graph, inputs=inputs)

File ~/miniconda3/envs/ML/lib/python3.8/site-packages/torch/autograd/__init__.py:173, in backward(tensors, grad_tensors, retain_graph, create_graph, grad_variables, inputs)
    168     retain_graph = create_graph
    170 # The reason we repeat same the comment below is that
    171 # some Python versions print out the first line of a multi-line function
    172 # calls in the traceback and some print out the last line
--> 173 Variable_execution_engine_run_backward( # Calls into the C++ engine to run the backward pass
    174     tensors, grad_tensors, retain_graph, create_graph, inputs,
    175     allow_unreachable=True, accumulate_grad=True)

RuntimeError: element 0 of tensors does not require grad and does not have a grad_fn
```

### 正确例子：建立一个变量，然后把值填进去

```
In [30]: a

Parameter containing:
        tensor([0.0592, 0.3881, 0.0161, 0.5500, 0.5207], requires_grad=True)

In [32]: A2 = torch.zeros([5])

In [33]: A2[0] = a[0] * 2
        A2[1] = torch.log(a[1]*1)
        A2[2] = a[2] ** 2
        A2[3] = torch.nn.functional.softplus(a[3] * 2)
        A2[4] = a[4]
        A2

Out[33]: tensor([1.1843e-01, 3.0607e-01, 2.5946e-04, 1.3874e+00, 5.2072e-01],
        grad_fn=<CopySlices>)

就可以使用了

In [34]: A2.sum().backward()

In [35]: a.grad

Out[35]: tensor([2.0000, 0.7363, 0.0322, 1.5006, 1.0000])
```

### 正确例子2：用cat或stack拼起来

```
In [62]: A3 = torch.stack([a[0]*2,
                         torch.log(a[1]*1),
                         a[2]**2,
                         torch.nn.functional.softplus(a[3]*2),
                         a[4]])

In [63]: a.grad.zero_()
        A3.sum().backward()

In [64]: a.grad

Out[64]: tensor([2.0000, 0.7363, 0.0322, 1.5006, 1.0000])
```

## 可学习参数有相互组合时，不能填进去，只能cat/stack

在某些情况下，神经网络的可学习参数需要经过一些变化，然后继续计算。例如在某些问题中，可学习参数 $R_1, R_2$ 需要经过一个黑箱过程才能变成 $\eta$ ，然后 $\eta$ 会参与后续的运算。也就是

$$[R_1, R_2] \rightarrow Black\ Box \rightarrow \eta \rightarrow NeuralNetwork$$

这个黑箱转换过程可以通过神经网络来实现，也就是

$$[R_1, R_2] \rightarrow NN \rightarrow \eta \rightarrow NeuralNetwork$$

有的时候 $R_1$ 和 $R_2$ 的比值是一个重要的因素。但是这个比值会随着normalization而被削弱。在这种情况下，人们往往人为添加一个特征，也就是训练一个3输入1输出的NN来拟合黑箱，也就是输入变成 $R_1, R_2, k$ ，其中 $k = \frac{R_1}{R_2}$ ，然后整个输入会被神经网络输入成为 $R_1^N, R_2^N, k^N$ ，注意 $k^N \neq \frac{R_1^N}{R_2^N}$ 。

然而，我们明显能看出 $R_1^N, R_2^N, k^N$ 只有2个自由度，也就是可学习参数最多只能有2个，否则就会产生矛盾。那么比如如果我们让可学习参数为 $R_1^N$ 和 $R_2^N$ ，然后推算出 $k^N = \mathcal{N}\left\{\frac{D(R_1^N)}{D(R_2^N)}\right\}$ ，其中 $\mathcal{N}(\cdot)$ 和 $D(\cdot)$ 表示normalization和denormalization的操作。然后把 $R_1^N$ 和 $R_2^N$ 推算出来的 $k^N$ 拼起来，放到NN里面计算 $\eta$ 。然而，这个过程会出现问题：

```
In [69]: torch.manual_seed(0)
        class Example(torch.nn.Module):
            def __init__(self):
                super().__init__()
                # learnable R1n and R2n
                self.paramm = torch.nn.Parameter(torch.rand(2),requires_grad=True)
                # exemplary NN convert [R1n, R2n, kn] to eta
                self.eta_estimator = torch.nn.Sequential(torch.nn.Linear(3,1))
                self.eta_estimator.train(False)
                # exemplary max and min of [R1, R2, k] for normalization and denormalization
                self.param_max = torch.tensor([2, 4, 2])
                self.param_min = torch.tensor([1, 2, 0])

            @property
            def Param(self):
                # calculate normalized R1
                paramm = torch.zeros([3])
                paramm[0] = self.paramm[0] # R1n
                paramm[1] = self.paramm[1] # R2n
                # denormalization
                param = paramm * (self.param_max - self.param_min) + self.param_min
                # calculate k=R1/R2
                param[2] = (param[0] / param[1]).detach()
                # normalization
                paramm = (param - self.param_min) / (self.param_max - self.param_min)
                return paramm

            @property
            def eta(self):
                return self.eta_estimator(self.Param)
```

可以看出上述代码以 $R_1^N$ 和 $R_2^N$ 为可学习参数，然后先经过denormalization求出真实值 $R_1$ 和 $R_2$ ，然后在利用真实值求出 $k$ ，再把真实的 $[R_1, R_2, k]$ 进行normalization，就可以用于放到NN里用于计算了：

```
In [70]: test = Example()
        print('可学习参数:')
        print(test.paramm)
        print('经过处理后的可学习参数')
        print(test.Param)
        print('利用可学习参数转换成的eta')
        print(test.eta)

可学习参数:
Parameter containing:
        tensor([0.4963, 0.7682], requires_grad=True)
经过处理后的可学习参数
        tensor([0.4963, 0.7682, 0.2115], grad_fn=<DivBackward0>)
利用可学习参数转换成的eta
        tensor([-0.4544], grad_fn=<AddBackward0>)

然而，由于 paramm 里的第三个参数k是前两个的组合，反向传递时会出现问题：

In [71]: test.Param.sum().backward()

RuntimeError                                Traceback (most recent call last)
Input In [71], in <cell line: 1>()
----> 1 test.Param.sum().backward()

File ~/miniconda3/envs/ML/lib/python3.8/site-packages/torch/_tensor.py:363, in Tensor.backward(self, gradient, retain_graph, create_graph, inputs)
    354 if has_torch_function_unary(self):
    355     return handle_torch_function(
    356         Tensor.backward,
    357         (self,),
    358         (...)
    361         create_graph=create_graph,
    362         inputs=inputs)
--> 363 torch.autograd.backward(self, gradient, retain_graph, create_graph, inputs=inputs)

File ~/miniconda3/envs/ML/lib/python3.8/site-packages/torch/autograd/__init__.py:173, in backward(tensors, grad_tensors, retain_graph, create_graph, grad_variables, inputs)
    168     retain_graph = create_graph
    170 # The reason we repeat same the comment below is that
    171 # some Python versions print out the first line of a multi-line function
    172 # calls in the traceback and some print out the last line
--> 173 Variable_execution_engine_run_backward( # Calls into the C++ engine to run the backward pass
    174     tensors, grad_tensors, retain_graph, create_graph, inputs,
    175     allow_unreachable=True, accumulate_grad=True)

RuntimeError: one of the variables needed for gradient computation has been modified by an inplace operation: [torch.FloatTensor [1]], which is output 0 of AsStridedBackward0, is at version 1; expected version 0 instead. Hint: enable anomaly detection to find the operation that failed to compute its gradient, with torch.autograd.set_detect_anomaly(True).
```

问题就在于 param 的第三个参数k是用 param[2]=param[0]/param[1] 填进去的。应该用cat（高维度）或者stack（0维）拼起来。

### 正确例子：

```
In [72]: torch.manual_seed(0)
        class Example(torch.nn.Module):
            def __init__(self):
                super().__init__()
                # learnable R1n and R2n
                self.paramm = torch.nn.Parameter(torch.rand(2),requires_grad=True)
                # exemplary NN convert [R1n, R2n, kn] to eta
                self.eta_estimator = torch.nn.Sequential(torch.nn.Linear(3,1))
                self.eta_estimator.train(False)
                # exemplary max and min of [R1, R2, k] for normalization and denormalization
                self.param_max = torch.tensor([2, 4, 2])
                self.param_min = torch.tensor([1, 2, 0])

            @property
            def Param(self):
                # calculate normalized R1
                paramm = torch.zeros([3])
                paramm[0] = self.paramm[0] # R1n
                paramm[1] = self.paramm[1] # R2n
                # denormalization
                param = paramm * (self.param_max - self.param_min) + self.param_min
                # calculate k=R1/R2
                param[2] = (param[0] / param[1]).detach()
                # normalization
                paramm = (param - self.param_min) / (self.param_max - self.param_min)
                return paramm

            @property
            def eta(self):
                return self.eta_estimator(self.Param)
```

```
In [73]: test = Example()
        print('可学习参数:')
        print(test.paramm)
        print('经过处理后的可学习参数')
        print(test.Param)
        print('利用可学习参数转换成的eta')
        print(test.eta)

可学习参数:
Parameter containing:
        tensor([0.4963, 0.7682], requires_grad=True)
经过处理后的可学习参数
        tensor([0.4963, 0.7682, 0.2115], grad_fn=<DivBackward0>)
利用可学习参数转换成的eta
        tensor([-0.4544], grad_fn=<AddBackward0>)

In [74]: test.Param.sum().backward()
        test.paramm.grad

Out[74]: tensor([1.1414, 0.8804])
```

## 其他解决方案

由于关键原因在 paramm 的第三个元素是前两个元素的组合，这会使PyTorch出现问题。我们可以采取另外的解决方案：既然只有2个自由度可以学习，可以将任意一个变量detach掉，这样也可以解决此问题，但是会牺牲掉detach掉的分支的信息，因此不推荐：

```
In [78]: torch.manual_seed(0)
        class Example(torch.nn.Module):
            def __init__(self):
                super().__init__()
                # learnable R1n and R2n
                self.paramm = torch.nn.Parameter(torch.rand(2),requires_grad=True)
                # exemplary NN convert [R1n, R2n, kn] to eta
                self.eta_estimator = torch.nn.Sequential(torch.nn.Linear(3,1))
                self.eta_estimator.train(False)
                # exemplary max and min of [R1, R2, k] for normalization and denormalization
                self.param_max = torch.tensor([2, 4, 2])
                self.param_min = torch.tensor([1, 2, 0])

            @property
            def Param(self):
                # calculate normalized R1
                paramm = torch.zeros([3])
                paramm[0] = self.paramm[0] # R1n
                paramm[1] = self.paramm[1] # R2n
                # denormalization
                param = paramm * (self.param_max - self.param_min) + self.param_min
                # calculate k=R1/R2
                param[2] = (param[0] / param[1]).detach()
                # normalization
                paramm = (param - self.param_min) / (self.param_max - self.param_min)
                return paramm

            @property
            def eta(self):
                return self.eta_estimator(self.Param)
```

```
In [79]: test = Example()
        print('可学习参数:')
        print(test.paramm)
        print('经过处理后的可学习参数')
        print(test.Param)
        print('利用可学习参数转换成的eta')
        print(test.eta)

可学习参数:
Parameter containing:
        tensor([0.4963, 0.7682], requires_grad=True)
经过处理后的可学习参数
        tensor([0.4963, 0.7682, 0.2115], grad_fn=<DivBackward0>)
利用可学习参数转换成的eta
        tensor([-0.4544], grad_fn=<AddBackward0>)

In [77]: test.Param.sum().backward()
        test.paramm.grad

Out[77]: tensor([1., 0.1])
```

## 总结

可以看出三个不同实验的结果都是一样的，但是第一个实验无法反向传播，后面两个可以。但是相比于第二个，第三个会把detach掉的分支的信息给忽略，这不是最优的选择。

```
In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:
```